

Robotics II: Control, Modeling and Learning	
Laboratory 6	
By Zarema Balgabekova	
In [278]:	<pre>import numpy as np import pandas as pd from sklearn.model_selection import train_test_split import keras from keras.models import Sequential from keras.layers import Dense, Flatten, Dropout from tensorflow.keras.optimizers import Adam from keras.utils import np_utils from keras import backend as K</pre>
The dataset of 5000 points was collected using dataset.py.	
In [279]:	<pre>#dict.csv contains 5000 data points data = pd.read_csv("dict2.csv", header = None, names = ["Angles", "Xy"])</pre>
In [280]:	<pre>train = data["Angles"].to_numpy() labels = data["Xy"].to_numpy()</pre>
In [282]:	<pre>X = list() Y = list() for i in range(len(train)):     labels[i] = labels[i].replace(' ', '')     labels[i] = labels[i].replace(' ', '')     labels[i] = labels[i].strip(['\n']).strip(' ')     train[i] = train[i].strip(['\n']).strip(' ')     result = [float(val) for val in train[i].split(',')]     X.append(result)     result = [float(val) for val in labels[i].split(' ')]     Y.append(result)</pre>
In [283]:	<pre>def rmse(y_true, y_pred):     return K.sqrt(K.mean(K.square(y_pred - y_true)))</pre>

## Trying different regression losses

### 1 RMSE (Root mean squared error)

In [284]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss=rmse, optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("RMSE: %.6f" % (scores))</pre>
RMSE: 0.062683	
2 MSE (Mean squared error)	

In [285]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSE: %.6f" % (scores))</pre>
MSE: 0.004353	
3 MAE (Mean absolute error)	

In [286]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_absolute_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MAE: %.6f" % (scores))</pre>
MAE: 0.052243	
4 MSLE (Mean squared logarithmic error)	

In [287]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000465	
5 Huber loss	

In [288]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='huber', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("Huber: %.6f" % (scores))</pre>
Huber: 0.001976	
6 Log cosh	

In [289]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='log_cosh', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("Log cosh: %.6f" % (scores))</pre>
Log cosh: 0.001451	

Therefore, MSLE loss should be used.

## Trying different optimizers

1 Adam	
For Adam optimizer, MSLE = 0.000465.	

In [290]:	<pre>from tensorflow.keras.optimizers import SGD from tensorflow.keras.optimizers import Adadelta from tensorflow.keras.optimizers import RMSprop from tensorflow.keras.optimizers import Adagrad from tensorflow.keras.optimizers import Adamax from tensorflow.keras.optimizers import Nadam from tensorflow.keras.optimizers import Ftrl</pre>
2 SGD	

In [291]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=SGD(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.050566	
3 Adadelta	

In [292]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adadelta(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.744811	
4 RMSprop	

In [293]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=RMSprop(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.001378	
5 Adagrad	

In [294]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adagrad(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.005846	
6 Adamax	

In [295]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adamax(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000792	
7 Nadam	

In [296]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Nadam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000735	
8 Ftrl	

In [297]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Ftrl(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.002130	

Therefore, Adam optimizer should be used. However, Nadam and Adamax also perform well.

## Changing layers

### 1 Changing the number of neurons for the first hidden layer

In [298]:	<pre>model = Sequential() model.add(Dense(32, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000516	

In [299]:	<pre>model = Sequential() model.add(Dense(64, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000842	

In [300]:	<pre>model = Sequential() model.add(Dense(5, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.001028	

We can see that the result does not improve. That's why we will keep 10 neurons for the first layer.

### 2 Adding more hidden layers

In [308]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(8, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000399	

In [307]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(4, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000543	

In [310]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(16, activation = 'relu')) model.add(Dense(8, activation = 'relu')) model.add(Dense(4, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000472	

It is seen that MSLE is the smallest when we add one more hidden layer with 8 neurons.

### 3 Decreasing the number of hidden layers

In [315]:	<pre>model = Sequential() model.add(Dense(10, input_dim =5, activation = 'relu')) model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.000766	

In [316]:	<pre>model = Sequential() model.add(Dense(2, activation='linear')) model.compile(loss='mean_squared_logarithmic_error', optimizer=Adam(0.01))  X_train, X_test, y_train, y_test = train_test_split(np.asarray(X), np.asarray(Y), test_size=0.20) y_train = np.delete(y_train, 2, 1) y_test = np.delete(y_test, 2, 1)  model.fit(X_train, y_train, epochs = 15, verbose = 0) scores = model.evaluate(X_test, y_test, verbose = 0) print("MSLE: %.6f" % (scores))</pre>
MSLE: 0.005244	

Decreasing the number of layers does not improve the result.

## Results

By using the provided parameters (neural network with two hidden layers, RMSE loss, and Adam optimizer), we obtained the following loss: **RMSE = 0.062683**. It was discovered that MSLE significantly decreases the loss. Trying different optimizers, we came to conclusion that Adam optimizer gives the best performance. Furthermore, the addition of one more hidden layer also improved the result. The best loss obtained is **MSLE = 0.000399**. Therefore, the result was improved by 157 times.